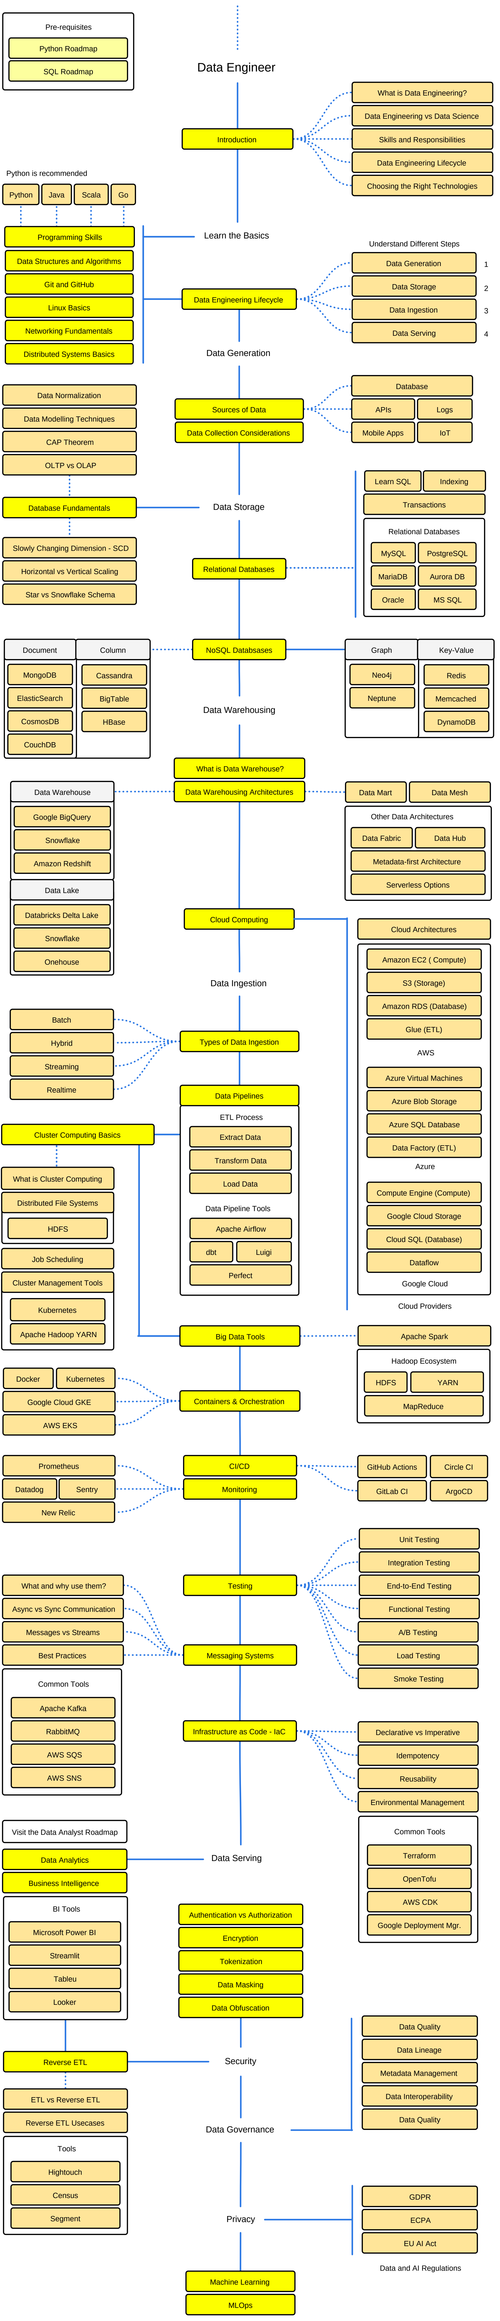
Pre-requisites
Python Roadmap
SQL Roadmap
Data Engineer
Introduction
What is Data Engineering?
Data Engineering vs Data Science
Skills and Responsibilities
Data Engineering Lifecycle
Choosing the Right Technologies
Python is recommended
Python
Java
Scala
Go
Programming Skills
Data Structures and Algorithms
Git and GitHub
Linux Basics
Networking Fundamentals
Distributed Systems Basics
Learn the Basics
Understand Different Steps
Data Engineering Lifecycle
Data Generation
Data Storage
Data Ingestion
Data Serving
1
2
3
4
Data Generation
Data Normalization
Data Modelling Techniques
CAP Theorem
OLTP vs OLAP
Sources of Data
Data Collection Considerations
Database
APIs
Logs
Mobile Apps
IoT
Database Fundamentals
Data Storage
Learn SQL
Indexing
Transactions
Relational Databases
MySQL
PostgreSQL
MariaDB
Aurora DB
Oracle
MS SQL
Slowly Changing Dimension - SCD
Horizontal vs Vertical Scaling
Star vs Snowflake Schema
Relational Databases
Document
Column
MongoDB
ElasticSearch
CosmosDB
CouchDB
Cassandra
BigTable
HBase
NoSQL Databsases
Graph
Key-Value
Neo4j
Neptune
Redis
Memcached
DynamoDB
Data Warehousing
What is Data Warehouse?
Data Warehouse
Google BigQuery
Snowflake
Amazon Redshift
Data Lake
Databricks Delta Lake
Snowflake
Onehouse
Data Warehousing Architectures
Data Mart
Data Mesh
Other Data Architectures
Data Fabric
Data Hub
Metadata-first Architecture
Serverless Options
Cloud Computing
Cloud Architectures
Amazon EC2 (Compute)
S3 (Storage)
Amazon RDS (Database)
Glue (ETL)
AWS
Azure Virtual Machines
Azure Blob Storage
Azure SQL Database
Data Factory (ETL)
Azure
Compute Engine (Compute)
Google Cloud Storage
Cloud SQL (Database)
Dataflow
Google Cloud
Cloud Providers
Data Ingestion
Batch
Hybrid
Streaming
Realtime
Types of Data Ingestion
Data Pipelines
ETL Process
Extract Data
Transform Data
Load Data
Data Pipeline Tools
Apache Airflow
dbt
Luigi
Perfect
Cluster Computing Basics
What is Cluster Computing
Distributed File Systems
HDFS
Job Scheduling
Cluster Management Tools
Kubernetes
Apache Hadoop YARN
Big Data Tools
Apache Spark
Hadoop Ecosystem
HDFS
YARN
MapReduce
Docker
Kubernetes
Google Cloud GKE
AWS EKS
Containers & Orchestration
Prometheus
Datadog
Sentry
New Relic
CI/CD
Monitoring
GitHub Actions
Circle CI
GitLab CI
ArgoCD
Testing
Unit Testing
Integration Testing
End-to-End Testing
Functional Testing
A/B Testing
Load Testing
Smoke Testing
What and why use them?
Async vs Sync Communication
Messages vs Streams
Best Practices
Messaging Systems
Common Tools
Apache Kafka
RabbitMQ
AWS SQS
AWS SNS
Infrastructure as Code - IaC
Declarative vs Imperative
Idempotency
Reusability
Environmental Management
Common Tools
Terraform
OpenTofu
AWS CDK
Google Deployment Mgr.
Visit the Data Analyst Roadmap
Data Analytics
Business Intelligence
Data Serving
BI Tools
Microsoft Power BI
Streamlit
Tableu
Looker
Authentication vs Authorization
Encryption
Tokenization
Data Masking
Data Obfuscation
Security
Reverse ETL
ETL vs Reverse ETL
Reverse ETL Usecases
Tools
Hightouch
Census
Segment
Data Governance
Data Quality
Data Lineage
Metadata Management
Data Interoperability
Data Quality
Privacy
GDPR
ECPA
EU AI Act
Data and AI Regulations
Machine Learning
MLOps